Conformal: an R package to calculate prediction errors in the conformal prediction framework

October 23, 2014

`conformal`	*conformal: an R package to calculate prediction errors in the conformal prediction framework*

Description

conformal permits the calculation of prediction errors in the conformal prediction framework: (i) p.values for classification, and (ii) confidence intervals for regression. The package is coded using R reference classes (OOP).

Details

Assessing the reliability of individual predictions is foremost in machine learning to determine the applicability domain of a predictive model, be it in the context of classification or regression. The applicability domain is usually defined as the amount (or the regions) of descriptor space to which a model can be reliably applied. Conformal prediction is an algorithm-independent technique, i.e. it works with any predictive method such as Support Vector Machines or Random Forests (RF), which outputs confidence regions for individual predictions in the case of regression, and p.values for categories in a classification setting [1,2].

Regression

In the conformal prediction framework [1,2], the datapoints in the training set are used to define how unlikely a new datapoint is with respect to the data presented to the model in the training phase. The *unlikeliness* (conformity) for a given datapoint, x, with respect to the training set is quantified with a nonconformity score, α, calculated with a nonconformity measure (e.g. `StandardMeasure`) [2], which here we define as:

$$\alpha = \frac{|y - \widetilde{y}|}{\widetilde{\rho}}$$

where α is the nonconformity score, y and $\widetilde{y}$ are respectively the observed and the predicted value calculated with an e point prediction model, and $\widetilde{\rho}$ is the predicted error for x calculated with an error model.
In order to calculate confidence intervals, we need a point prediction model, to predict the response

1

variable (y), and an error model, to predict errors in prediction ($\widetilde{\rho}$). The point prediction and error models can be generated with any machine learning algorithm. Both the point prediction and error models need to be trained with cross-validation in order to calculate the vector of nonconformity scores for the training set, $D_i = \{x_i\}_i^{N_{tr}}$.

The cross-validation predictions generated when training the point prediction model serve to calculate the errors in prediction for the datapoints in the training set, $y_i - \widetilde{y}_i$. The error model is then generated by training a machine learning model on the training set using these errors as the dependent variable. The (i) cross-validated predictions from the point prediction model, and (ii) the cross-validated errors in prediction from the error model, are used to generate the vector of nonconformity scores for the training set. This vector, after being sorted in increasing order, can be defined as:

$$\alpha_{tr} = \{\alpha_{tr\ i}\}_i^{N_{tr}}$$

where N_{tr} is the number of datapoints in the training set.

To generate the confidence intervals for an external set, $D_{ext} = \{x_{ext}\}_j^{N_{ext}}$, we have to define a confidence level, ϵ. The α value associated to the user-defined confidence level, α_ϵ, is calculated as:

$$\alpha_\epsilon = \alpha_{tr\ i}\ \ if\ \ i \equiv |N_{tr} * \epsilon|$$

where $\equiv$ indicates equality. Next, the errors in prediction, $\widetilde{\rho}_{ext}$, and the value for the response variable, y_{ext}, for the datapoints in the external dataset are predicted with the error and the point prediction models, respectively.

Individual confidence intervals (CI) for each datapoint in the external set are derived from:

$$CI_{ext\ j} = |y_{ext\ j} - \widetilde{y}_{ext\ j}| = \alpha_\epsilon * \widetilde{\rho}_{ext\ j}$$

The confidence region (CR) is finally defined as:

$$CR = \widetilde{y}_{ext\text{æ}} + / - CI_{ext\text{æ}}$$

The interpretation of the confidence regions is straightforward. For instance, it we choose a confidence level of 80% the true value for new datapoints will lie outside the predicted confidence regions in at most 20% of the cases.

Classification

Initially, a Random Forest classifier is trained on the training set using k-fold cross-validation.

In the case of classification, the nonconformity scores are calculated on a *per* class basis. These are calculated as the ratio between the number of trees in the forest voting for a given class divided by the total number of trees (label-wise Mondrian off-line inductive conformal prediction -MICP-) [3]. For instance, in a binary classification example, if 87 trees from a Random Forest model comprising 100 trees classify a datapoint as belonging to class A, the nonconformity score (or probability) for this class would be 0.87 (87%), whereas its value for class B would be 0.13. This process generates a matrix (nonconformity scores matrix) which rows correspond to the datapoints in the training set, and its columns to the number of distinct classes (two in the binary classification example) (Figure 1A). Here, we have implemented the pipeline proposed by Norinder et al. 2014 [2] using Ranfom Forest models. Nevertheless, other ensemble methods could be used to calculate the nonconformity scores.
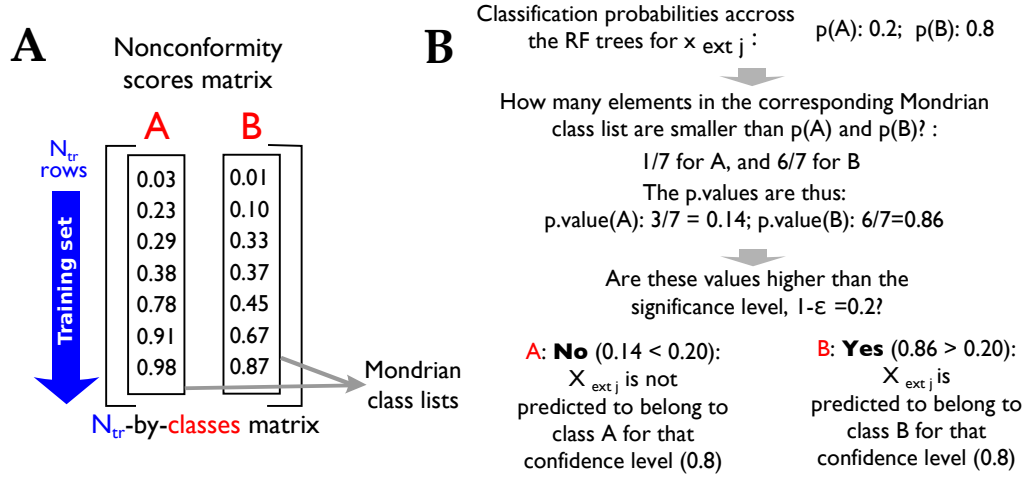
Figure 1. Calculation of conformal prediction errors (p.values) in a binary classification example considering a confidence level of 0.80.

Next, each column of the matrix is sorted in increasing order. These columns are called Mondrian class lists (MCL) (Figure 1A). As in regression, a confidence level, ϵ, needs to be specified. We define significance as $1 - \epsilon$.

The model trained on the whole dataset is used to classify the datapoints comprised in the external dataset (Figure 1B). Let's consider one datapoint from the external set, namely $x_{ext\ j}$. The number of trees in the Random Forest voting for each class is computed, which enables the caculation of the nonconformity scores or probabilities (p) for that point, $x_{ext\ j}$. In the binary case, this is defined as:

$$p(x_{ext\ j}; A) = \frac{N_{trees}votingA}{N_{trees}}; \quad p(x_{ext\ j}; B) = \frac{N_{trees}votingB}{N_{trees}}$$

To calculate the p.values for each class, the number of elements in the corresponding Mondrian class list smaller than the probability values, i.e. $p(x_{ext\ j}; A)$ and $p(x_{ext\ j}; B)$, is divided by the number of datapoints in the training set, N_{tr}:

$$p.value(x_{ext\ j}; A) = \frac{|\{MCL(A) < P(x_{ext\ j}; A)\}|}{N_{tr}}; \quad p.value(x_{ext\ j}; B) = \frac{|\{MCL(B) < P(x_{ext\ j}; B)\}|}{N_{tr}}$$

Finally, these p.values are compared to the significance level defined by the user $(1 - \epsilon)$. For a datapoint to be predicted to belong to a given class, the p.value needs to be higher than the significance level. For instance, if $p.value(x_{ext\ j}; A) = 0.46$ and $p.value(x_{exti}; B) = 0.18$, with a significance level of 0.2, $x_{ext\ j}$ would be predicted to belong to class A, but not to B. If both $p.value(x_{ext\ j}; A)$ and $p.value(x_{ext\ j}; B)$ were higher than the significance level, $x_{ext\ j}$ would be predicted to belong to both classes. Similarly, if both p.values were smaller than the significance level, $x_{ext\ j}$ would be predicted to belong to neither class A nor to class B.

Author(s)

Isidro Cortes <isidrolauscher@gmail.com>. conformal: an R package to calculate prediction errors in the conformal prediction framework.

References

[1] Shafer et al. JMLR, 2008, 9, pp 371-421. http://machinelearning.wustl.edu/mlpapers/paper_files/shafer08a.pdf

[2] Norinder et al. J. Chem. Inf. Model., 2014, 54 (6), pp 1596-1603. DOI: 10.1021/ci5001168 http://pubs.acs.org/doi/abs/10.1021/ci5001168

[3] Dmitry Devetyarov and Ilia Nouretdinov, Artificial Intelligence Applications and Innovations, 2010, 339, pp 37-44. DOI: 10.1007/978-3-642-16239-8_8 http://link.springer.com/chapter/10.1007%2F978-3-642-16239-8_8#

ConformalClassification

Conformal Prediction For Classification

Description

R reference class to calculate p.values for individual predictions according to the conformal prediction framework.

Details

The reference class ConformalClassification contains the following fields:

- ClassificationModel: stores a classification Random Forest model.
- confidence: stores the user-defined confidence level.
- data.new: stores the descriptors corresponding to an external set.
- NonconformityScoresMatrix: stores the non conformity scores matrix.
- ClassPredictions: stores the class predictions calculated for the external set.
- p.values: a list storing
    - P.values: a data.frame containing the p.values calculated for the external set. Rows are indexed by datapoints, whereas columns are indexed by classes. The names of the rows correspond to the names of the rows in the external set.
    - Significance_p.values: a data.frame reporting the significance of the p.values (where 1 means significant, and 0 not significant), according to the user-defined confidence level, ϵ (the default value is 0.8). Rows are indexed by datapoints, whereas columns are indexed by classes. The names of the rows correspond to the names of the rows in the external set.

The class ConformalClassification contains the following methods:

- initialize: this method is called when you create an instance of the class. The default value for the confidence level is 0.8.
- CalculateCVScores: this method calculates the non conformitity scores (or probabilities) matrix from the cross-validation predictions of the input randomForest model (trained with k-fold cross-validation). The non conformity scores matrix is stored in the field NonconformityScoresMatrix.
- CalculatePValues: this method calculates the p.values for the datapoints in a external set. The class predictions are stored in the field ClassPredictions, whereas the p.values and their significance, according to the user defined confidence level, are stored in the field p.values.

Author(s)

Isidro Cortes-Ciriano <isidrolauscher@gmail.com>

References

Norinder et al. J. Chem. Inf. Model., 2014, 54 (6), pp 1596-1603 DOI: 10.1021/ci5001168 http://pubs.acs.org/doi/abs/10.1021/ci5001168

See Also

ConformalRegression

Examples

```
# Optional for parallel training
#library(doMC)
#registerDoMC(cores=4)

data(LogS)

# convert data to categorical
LogSTrain[LogSTrain > -4] <- 1
LogSTrain[LogSTrain <= -4] <- 2
LogSTest[LogSTest > -4] <- 1
LogSTest[LogSTest <= -4] <- 2

LogSTrain <- factor(LogSTrain)
LogSTest <- factor(LogSTest)

algorithm <- "rf"

trControl <- trainControl(method = "cv",  number=5,savePredictions=TRUE)
set.seed(3)
model <- train(LogSDescsTrain, LogSTrain,
        algorithm,type="classification",
        trControl=trControl,predict.all=TRUE,
        keep.forest=TRUE,norm.votes=TRUE,
        ntree=500)


# Instantiate the class and get the p.values
example <- ConformalClassification$new()
example$CalculateCVScores(model=model)
example$CalculatePValues(new.data=LogSDescsTest)
# we get the p.values:
example$p.values$P.values
# we get the significance of these p.values.
example$p.values$Significance_p.values
```

ConformalRegression *Conformal Prediction for Regression*

Description

R reference class to create and visualize confidence intervals for individual predictions according to the conformal prediction framework.

Details

The reference class ConformalRegression contains the following fields:

- PointPredictionModel: stores a point prediction model.
- ErrorModel: stores an error model. Error models can be trained with the function `ErrorModel`.
- confidence: stores the level of confidence used to calculate the confidence intervals. This value is defined when instantiating a new class. Values are in the 0-1 range. Interpretation: for instance, a confidence level of 0.8 (80%) means that, at most, 20% of the confidence intervals will not contain the true value.
- data.new: stores the descriptors corresponding to an external set.
- alphas: stores the nonconformity scores for the datapoints used to train the point prediction model (PointPredictionModel) with the method CalculateAlphas.
- errorPredictions: errors in prediction for an external set predicted with an error model (stored in the field ErrorModel). These errors are generated when calling the method GetConfidenceIntervals. These errors serve to calculate the nonconformity scores (alphas) for the an external set. In the conformal prediction framework, the confidence intervals are derived from the nonconformity values (see methods CalculateAlphas and GetConfidenceIntervals).
- pointPredictions: point predictions for an external set calculated with the point prediction model (stored in the field PointPredictionModel). These predictions are generated with the method GetConfidenceIntervals.
- intervals: numeric vector with the errors in prediction for the external set (data.new) calculated in the conformal prediction framework (not with an error model). These intervals are calculated when calling the method GetConfidenceIntervals.
- plot: stores a correlation plot for the observed against the predicted values, with individual confidence intervals, for the datapoints in an external set. The plot is a ggplot2 object which can be further customized. The plot is generated with the method CorrelationPlot.

The class ConformalRegression contains the following methods:

- initialize: this method is called when you create an instance of the class. The default value for the confidence level is 0.8.
- CalculateAlphas: this method calculates the vector of nonconformity scores for the datapoints in the traning set. These scores (or alphas) are stored in the field alphas. This method requires a point prediction model (argument model), an error model (argument error_model), and a nonconformity measure (argument ConformityMeasure), such as `StandardMeasure`.
- GetConfidenceIntervals: this methods calculates confidence intervals for individual predictions in the conformal prediction framework. The methods requires an external set for which the confidence intervals will be calculated. The dimensionality of these descriptors must be the same as the one used for the datapoints used to train the point prediction and the error model. The method uses the point prediciton and the error models stored in the fields PointPredictionModel and ErrorModel, respectively. Confidence intervals are calculated according to Norinder et al. 2014. Confidence intervals are defined as: point predictions (stored in the field pointPredictions) +/- the output of the method GetConfidenceIntervals, which is stored in the field intervals.
- CorrelationPlot: this method generates a correlation plot for the observed against the predicted values for an external set, along with individual confidence intervals. The plot is stored in the field plot.

Author(s)

Isidro Cortes-Ciriano <isidrolauscher@gmail.com>

References

Norinder et al. J. Chem. Inf. Model., 2014, 54 (6), pp 1596-1603 DOI: 10.1021/ci5001168 http://pubs.acs.org/doi/abs/10.1021/ci5001168

See Also

conformal

Examples

```
############################################
### Example
############################################

# Optional for parallel training
#library(doMC)
#registerDoMC(cores=4)

data(LogS)

algorithm <- "svmRadial"
tune.grid <- expand.grid(.sigma = expGrid(power.from=-10, power.to=-6, power.by=1, base=2),
                         .C = expGrid(power.from=4, power.to=10, power.by=2, base=2))
trControl <- trainControl(method = "cv",  number=5,savePredictions=TRUE)
set.seed(3)
model <- train(LogSDescsTrain, LogSTrain, algorithm,
               tuneGrid=tune.grid,
               trControl=trControl)


# Train an error model
error_model <- ErrorModel(PointPredictionModel=model,x.train=LogSDescsTrain,
                          savePredictions=TRUE,algorithm=algorithm,
                          trControl=trControl, tune.grid=tune.grid)

# Instantiate the class and get the confidence intervals
example <- ConformalRegression$new()
example$CalculateAlphas(model=model,error_model=error_model,ConformityMeasure=StandardMeasure)
example$GetConfidenceIntervals(new.data=LogSDescsTest)
example$CorrelationPlot(obs=LogSTest)
example$plot
```

ErrorModel   *Calculation of an error model*

Description

This function permits the calculation of an error model from (i) a training set, and (ii) a caret model trained on this set to predict the response variable of interest using cross-validation (point prediction model). The cross-validation predictions are extracted from the point prediction model. The errors

in prediction for the cross-validation predictions are then calculated. These errors which serve as the response variable for the error model (i.e. the error model predicts errors in prediction). The error model uses as descriptors the same descriptors used to train the point prediction model. These descriptors are input to the function ErrorModel through the argument "x.train".

Usage

```
ErrorModel(PointPredictionModel, x.train, algorithm = "svmRadial", ...)
```

Arguments

PointPredictionModel

> Point prediction model from which the cross-validation predictions will be extracted.

x.train   Descriptors for the datapoints in the training set used to train the point prediction model, and which will also serve to train the error model.

algorithm   The machine learning algorithm to be used to train the error model. The defaul value is Support Vector Machine with radial kernel ("svmRadial").

...   Additional arguments that can be passed to the train function from the R package caret to train the error model.

Value

A list of class train containing the error model (caret model).

Author(s)

Isidro Cortes-Ciriano <isidrolauscher@gmail.com>

References

http://caret.r-forge.r-project.org/training.html

Kuhn (2008), "Building Predictive Models in R Using the caret" (http://www.jstatsoft.org/v28/i05/)

See Also

[ConformalRegression](#) [GetCVPreds](#)

---

expGrid   *Exponential Grid Definition*

Description

The function defines an exponential series, which can be used, e.g. when defining the parameter space when training some models such as Support Vector Machines or Gaussian Processes.

Usage

```
expGrid(power.from, power.to, power.by, base)
```

Arguments

`power.from`	The starting exponential of the series.
`power.to`	The latest exponential of the series.
`power.by`	The exponential step of the series.
`base`	The base of the exponential series.

Value

A vector with the exponential series.

Author(s)

Isidro Cortes <isidrolauscher@gmail.com>

Examples

```
expGrid(power.from=-10,power.to=10,power.by=2,base=10)
```

`GetCVPreds`	*Extract the cross-validation predictions for the datapoints in the training set*

Description

This function extracts the cross-validation predictions from a caret model trained with cross-validation. If grid-search is used to optimize the hyperparamter values, the function GetCVPreds extracts the cross-validation predictions corresponding to the optimal hyperparameter values.

Usage

```
GetCVPreds(model)
```

Arguments

`model`	A caret model trained with cross-validation.

Value

A data frame containing the observed and the predicted values for the datapoints in the training set, their index according to the training set, the optimal hyperparameter values, and the fold index.

Author(s)

Isidro Cortes-Ciriano <isidrolauscher@gmail.com>

LogS   *Small Molecule Solubility (LogS) Data*

Usage

```
data(LogS)
```

Details

This dataset comprises the aqueous solubility (S) values at a temperature of 20-25 $^\circ$C in mol/L, expressed as logS, for 1,708 small molecules reported by Wang et al. Compound structures were standardized with the function StandardiseMolecules from the R package camb using the default parameters: (i) all molecules were kept irrespective of the numbers of fluorines, iodines, chlorine, and bromines present in their strucuture, or (ii) of their molecular mass. 905 one-dimensional topological and physicochemical descriptors were calculated with the function GeneratePadelDescriptors from the R package camb which invokes the PaDEL-Descriptor Java library. Near zero variance and highly-correlated descriptors were removed with the functions (i) RemoveNearZeroVarianceFeatures (cut-off value of 30/1), and (ii) RemoveHighlyCorrelatedFeatures (cut-off value of 0.95) After applying these steps the dataset consists of 1,606 molecules encoded with 211 descriptors.

Using data(LogS) exposes 4 objects:

(i) LogSDescsTrain is a data frame with PaDEL descriptors for the datapoints in the training set (70% of the data). (ii) LogSTrain is a numeric vector containing the data solubility values for the datapoints in the training set. (iii) LogSDescsTest is a data frame with PaDEL descriptors for the datapoints in the test set (30% of the data). (iv) LogSTest is a numeric vector containing the data solubility values for the datapoints in the test set.

References

Wang et al. J. Chem. Inf. Model., 2007, 47 (4), pp 1395-1404 DOI: 10.1021/ci700096r http://pubs.acs.org/doi/abs/10.1021/ci700096r

Examples

```
# To use the data
data(LogS)
```

StandardMeasure   *Default Nonconformity measure used in regression (ConformalRegression class)*

Description

Nonconformity measure used by default in the class ConformalRegression to calculate nonconformity scores (alpha). StandardMeasure implements the following nonconformity measure:

alpha = |obs - pred| / pred_error

where alpha is the nonconformity measure, obs is the observed value, pred the predicted value with a point prediction model, and pred_error is the error in prediction predicted with an error model. See conformal for further information about how to derive individual confidence intervals.

Usage

```
StandardMeasure(obs, pred, error)
```

Arguments

obs	Observed values for the response variable.
pred	Predicted values for the response variable.
error	Predicted errors calculated with an error model.

Value

A numeric vector with the nonconformity scores (alpha).

Author(s)

Isidro Cortes-Ciriano <isidrolauscher@gmail.com>

References

Norinder et al. J. Chem. Inf. Model., 2014, 54 (6), pp 1596-1603 DOI: 10.1021/ci5001168 http://pubs.acs.org/doi/abs/10.1021/ci5001168

See Also

ConformalRegression

Index